

# Storage Management

# Memory allocation techniques


- Uniprogramming
- Fixed-partition multiprogramming
- Variable-partition multiprogramming
- Paging
- Virtual memory

# Uniprogramming

- Operating system resides at one end of memory
- Application at the other end
- System only executes one job at time
- Example: Pep/8 operating system
- Disadvantages: Inflexible, CPU time wasted waiting for I/O

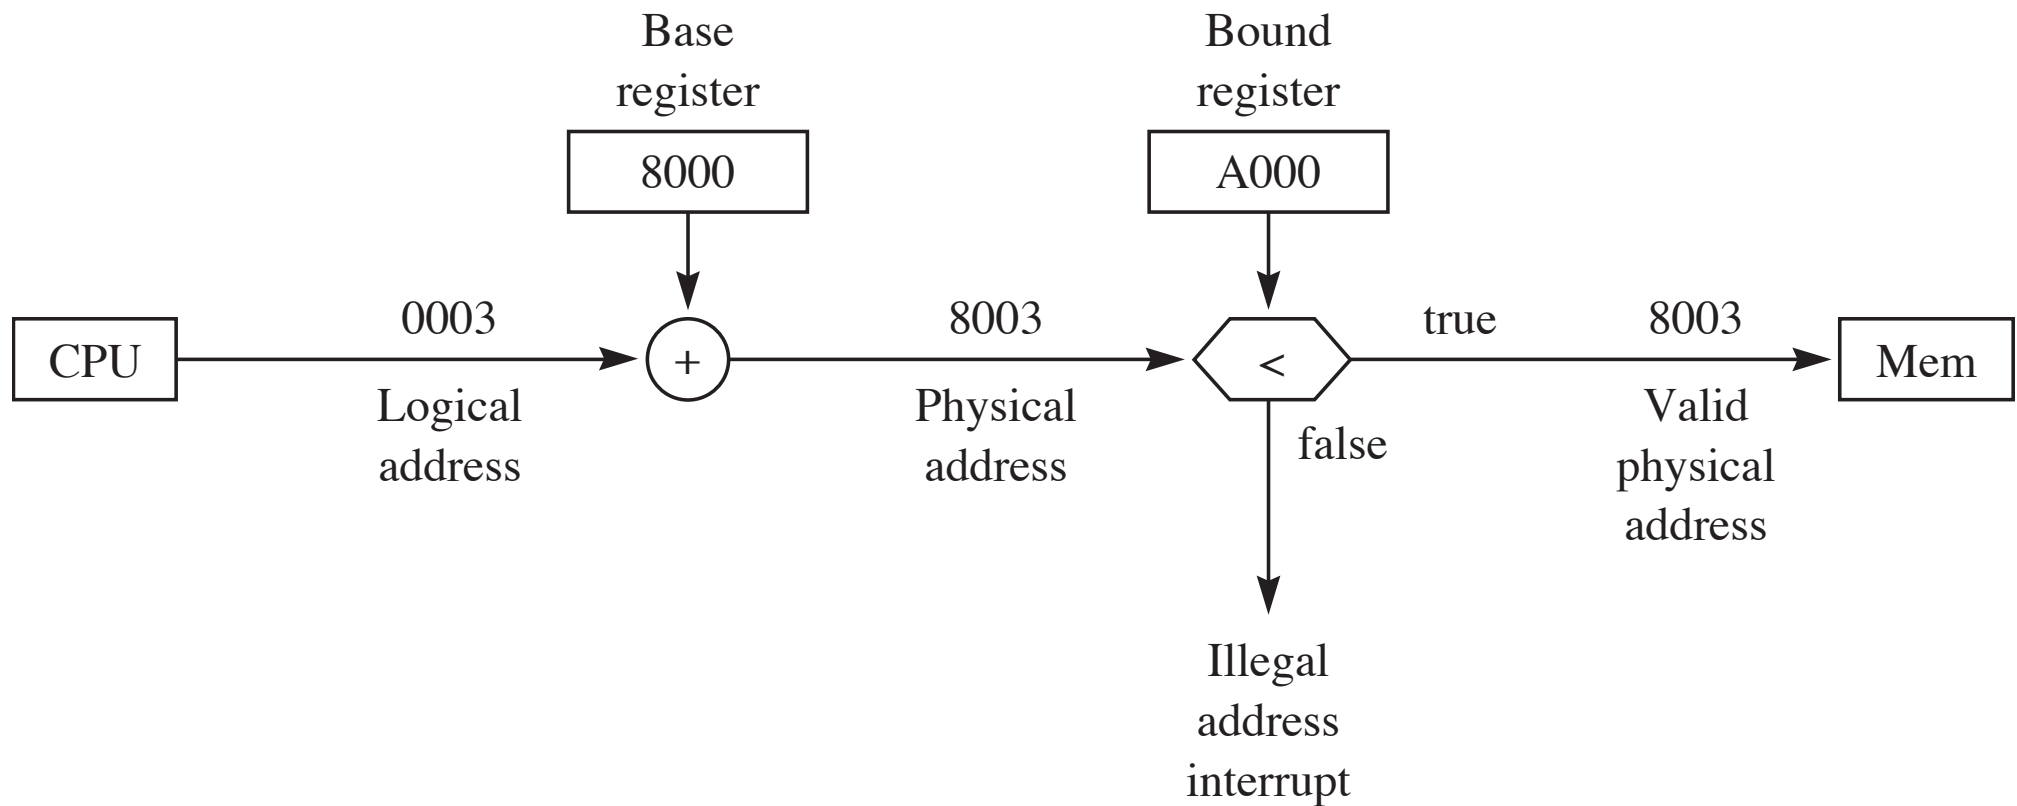
# Fixed-partition multiprogramming

- Operating system in one fixed reserved partition of memory
- Multiple processes in fixed partitions of memory
- Must solve the address problem



# Logical addresses

- Logical address is the address generated by the assembler assuming the program begins at address 0
- Physical address =  
$$\text{logical address} + \text{partition address}$$
- Base register converts from logical to physical
- Bound register keeps program isolated




# Problems with fixed-partitions

- Scheduling a small job in a large partition because the small partitions are all used
- Determining the optimal partition in the first place




# Variable-partition multiprogramming


- Establish a partition only when a job is loaded into memory
- The size of the partition can match the size of the job
- A region available for use by an incoming job is a *hole*




Job	Size	Action
J1	12	Start
J2	8	Start
J3	12	Start
J4	4	Start
J5	4	Start
J1	12	Stop
J6	4	Start
J5	4	Stop
J7	8	Start
J8	8	Start




(a) J1 to J5 start.




(b) J1 stops.




(c) J6 starts.




(d) J5 stops.




(e) J7 starts.




(a) Shifting all jobs to the top.




(b) Shifting only J6.




**(a)** J1 to J5 start.




**(b)** J1 stops.



**(c)** J6 starts.



**(d)** J5 stops.



**(e)** J7 starts.

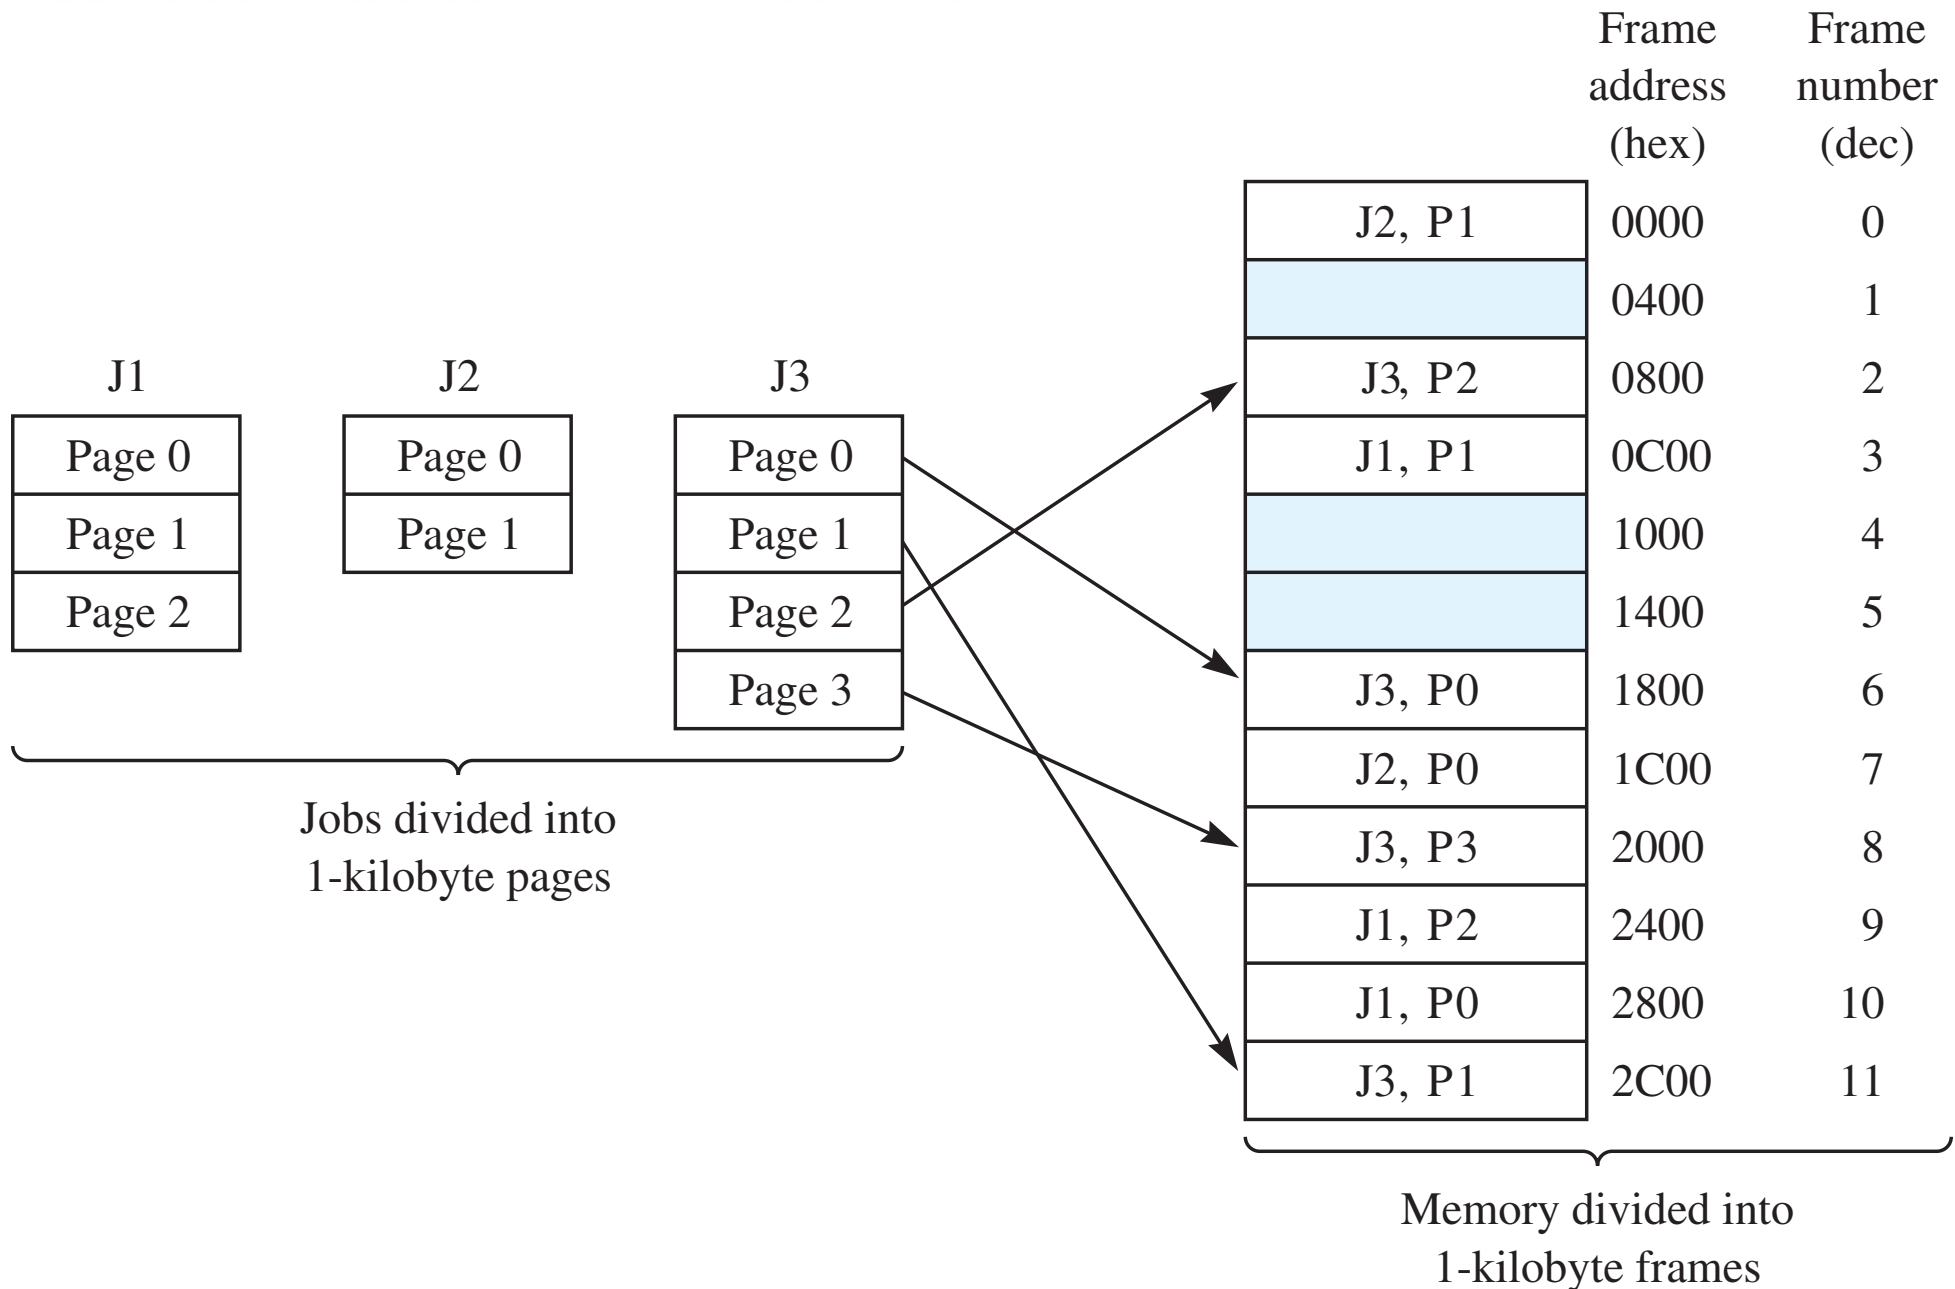
# Problems with variable partitions


- Fragmentation
- Consolidating holes is time-consuming

# Paging


- Rather than coalesce several small holes to fit the program, fragment the program to fit the holes
- A job is divided into pages
- Main memory is divided into frames, each one the same size as a page
- No coalescing of holes is ever required



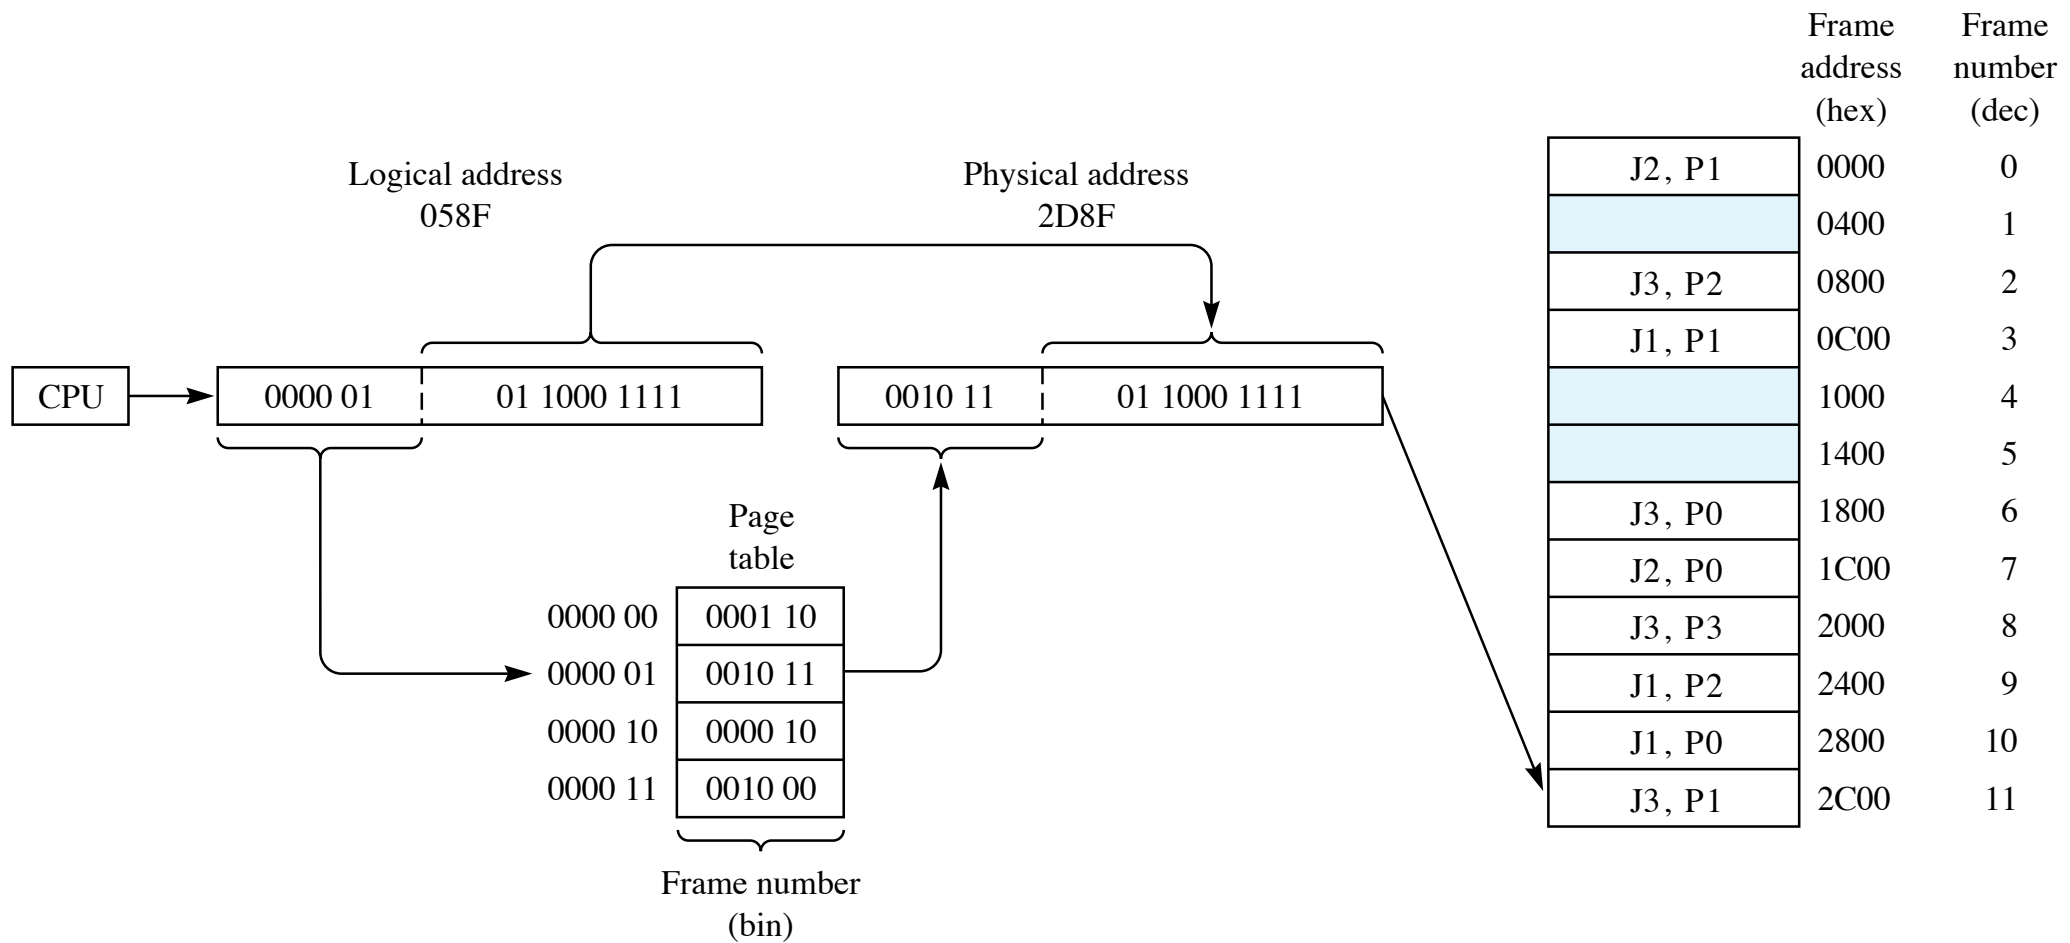





(a) Logical address.



(b) Physical address.



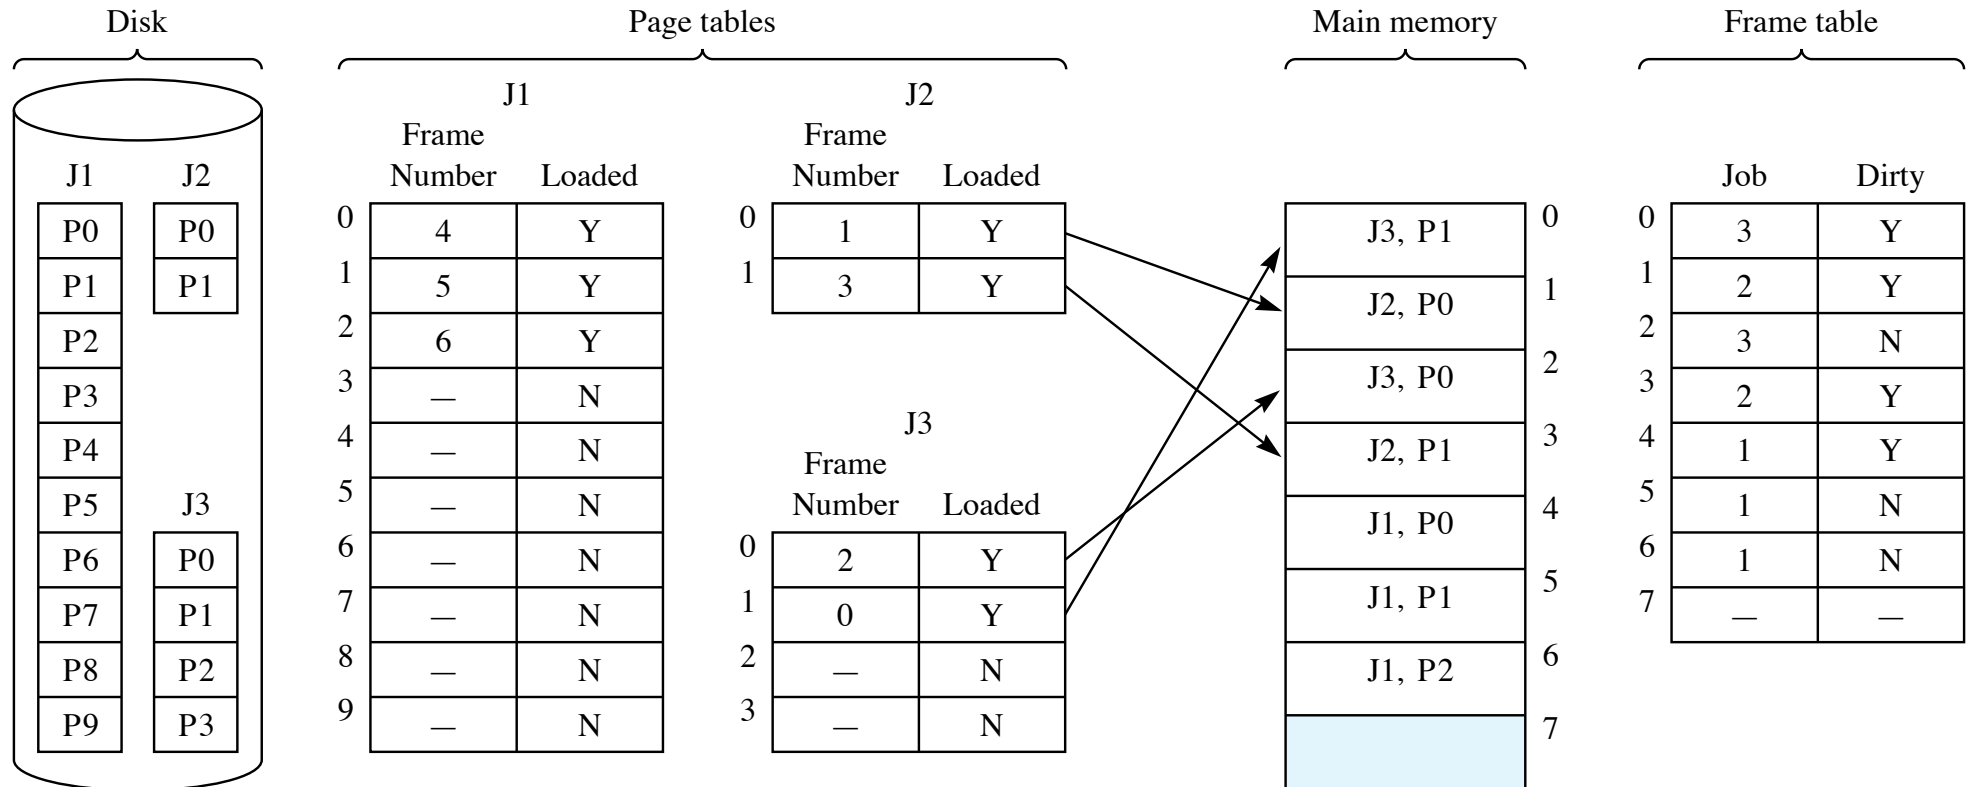


# Problem with paging

- To execute a program, the entire program must be loaded into memory
- Most large programs have many sections of code that never execute
- Memory is used inefficiently with parts of the program taking up main memory unnecessarily

# Virtual memory

- Cycle pages of the program from disk into memory only when they need to be executed
- The page that is executing in memory together with the pages in memory that have recently been executed is the program's *working set*
- As the program progresses, pages enter and leave the working set



# Page tables

- One page table for each program
- Converts logical address to physical address as in paging
- Loaded bit is 1 if the page is in memory
- A *page fault* occurs if the program needs to read or write a page that is not in memory
  - ▶ Page is loaded into an empty frame
  - ▶ If no empty frames, then a page is replaced



# Frame tables

- One frame table with an entry for each frame
- Dirty bit initialized to 0 when page is first loaded into memory
- Set to 1 on a `STr` to the frame, not on a `LDr` from the frame
- When a page is replaced it is written back to disk only if the dirty bit is set to 1

# Frame allocation

- Before a job starts executing, how many frames should be allocated for that job?
- System can allocate frames proportional to the physical size of the code

# Page replacement

- First in, first out (FIFO)
  - ▶ On a page fault, select the page to be replaced as the one that first entered the set of loaded pages
- Least recently used (LRU)
  - ▶ On a page fault, select the page to be replaced as the one that was least recently read from or written to

## FIFO


	6	8	3	8	6	0	3	6	3	5	3	6		Page references
—	6	8	3	3	3	0	0	6	6	5	3	3		
—	—	6	8	8	8	3	3	0	0	6	5	5		Loaded pages
—	—	—	6	6	6	8	8	3	3	0	6	6		
	F	F	F			F		F		F	F			Page fault

# FIFO


	6	8	3	8	6	0	3	6	3	5	3	6	Page references
	<div style="border: 1px solid black; padding: 2px;">—</div>	<div style="border: 1px solid black; padding: 2px;">6</div>	<div style="border: 1px solid black; padding: 2px;">8</div>	<div style="border: 1px solid black; padding: 2px;">3</div>	<div style="border: 1px solid black; padding: 2px;">3</div>	<div style="border: 1px solid black; padding: 2px;">3</div>	<div style="border: 1px solid black; padding: 2px;">0</div>	<div style="border: 1px solid black; padding: 2px;">0</div>	<div style="border: 1px solid black; padding: 2px;">0</div>	<div style="border: 1px solid black; padding: 2px;">0</div>	<div style="border: 1px solid black; padding: 2px;">5</div>	<div style="border: 1px solid black; padding: 2px;">5</div>	<div style="border: 1px solid black; padding: 2px;">6</div>
	<div style="border: 1px solid black; padding: 2px;">—</div>	<div style="border: 1px solid black; padding: 2px;">—</div>	<div style="border: 1px solid black; padding: 2px;">6</div>	<div style="border: 1px solid black; padding: 2px;">8</div>	<div style="border: 1px solid black; padding: 2px;">8</div>	<div style="border: 1px solid black; padding: 2px;">8</div>	<div style="border: 1px solid black; padding: 2px;">3</div>	<div style="border: 1px solid black; padding: 2px;">3</div>	<div style="border: 1px solid black; padding: 2px;">3</div>	<div style="border: 1px solid black; padding: 2px;">3</div>	<div style="border: 1px solid black; padding: 2px;">0</div>	<div style="border: 1px solid black; padding: 2px;">0</div>	<div style="border: 1px solid black; padding: 2px;">5</div>
	<div style="border: 1px solid black; padding: 2px;">—</div>	<div style="border: 1px solid black; padding: 2px;">—</div>	<div style="border: 1px solid black; padding: 2px;">—</div>	<div style="border: 1px solid black; padding: 2px;">6</div>	<div style="border: 1px solid black; padding: 2px;">6</div>	<div style="border: 1px solid black; padding: 2px;">6</div>	<div style="border: 1px solid black; padding: 2px;">8</div>	<div style="border: 1px solid black; padding: 2px;">8</div>	<div style="border: 1px solid black; padding: 2px;">8</div>	<div style="border: 1px solid black; padding: 2px;">8</div>	<div style="border: 1px solid black; padding: 2px;">3</div>	<div style="border: 1px solid black; padding: 2px;">3</div>	<div style="border: 1px solid black; padding: 2px;">0</div>
	<div style="border: 1px solid black; padding: 2px;">—</div>	<div style="border: 1px solid black; padding: 2px;">—</div>	<div style="border: 1px solid black; padding: 2px;">—</div>	<div style="border: 1px solid black; padding: 2px;">—</div>	<div style="border: 1px solid black; padding: 2px;">—</div>	<div style="border: 1px solid black; padding: 2px;">—</div>	<div style="border: 1px solid black; padding: 2px;">6</div>	<div style="border: 1px solid black; padding: 2px;">6</div>	<div style="border: 1px solid black; padding: 2px;">6</div>	<div style="border: 1px solid black; padding: 2px;">6</div>	<div style="border: 1px solid black; padding: 2px;">8</div>	<div style="border: 1px solid black; padding: 2px;">8</div>	<div style="border: 1px solid black; padding: 2px;">3</div>
	F	F	F			F				F		F	Page fault

# Belady's anomaly

- In general, the greater the number of frames allocated to a program, the fewer the number of page faults
- In a few cases with FIFO, an increase in the number of frames increases the number of page faults
- Example page reference sequence  
0, 1, 2, 3, 0, 1, 4, 0, 1, 2, 3, 4



**(a)** The expected effect of more frames on the number of page faults.



**(b)** Belady's anomaly with the FIFO replacement algorithm.


## LRU

	6	8	3	8	6	0	3	6	3	5	3	6		Page references
—	6	8	3	8	6	0	3	6	3	5	3	6		
—	—	6	8	3	8	6	0	3	6	3	5	3		Loaded pages
—	—	—	6	6	3	8	6	0	0	6	6	5		
	F	F	F			F	F			F				Page fault




# File management

- Create a new file
- Delete a file
- Rename a file
- Open a file for editing
- Read the next data item from the file



(a) A hard disk drive.



(b) A single disk.

# Contributions to the disk access time

- Seek time
  - ▶ Time for head to reach cylinder
- Latency
  - ▶ Time for start of sector to rotate to head
- Transmission time
  - ▶ Time for sector to pass under head

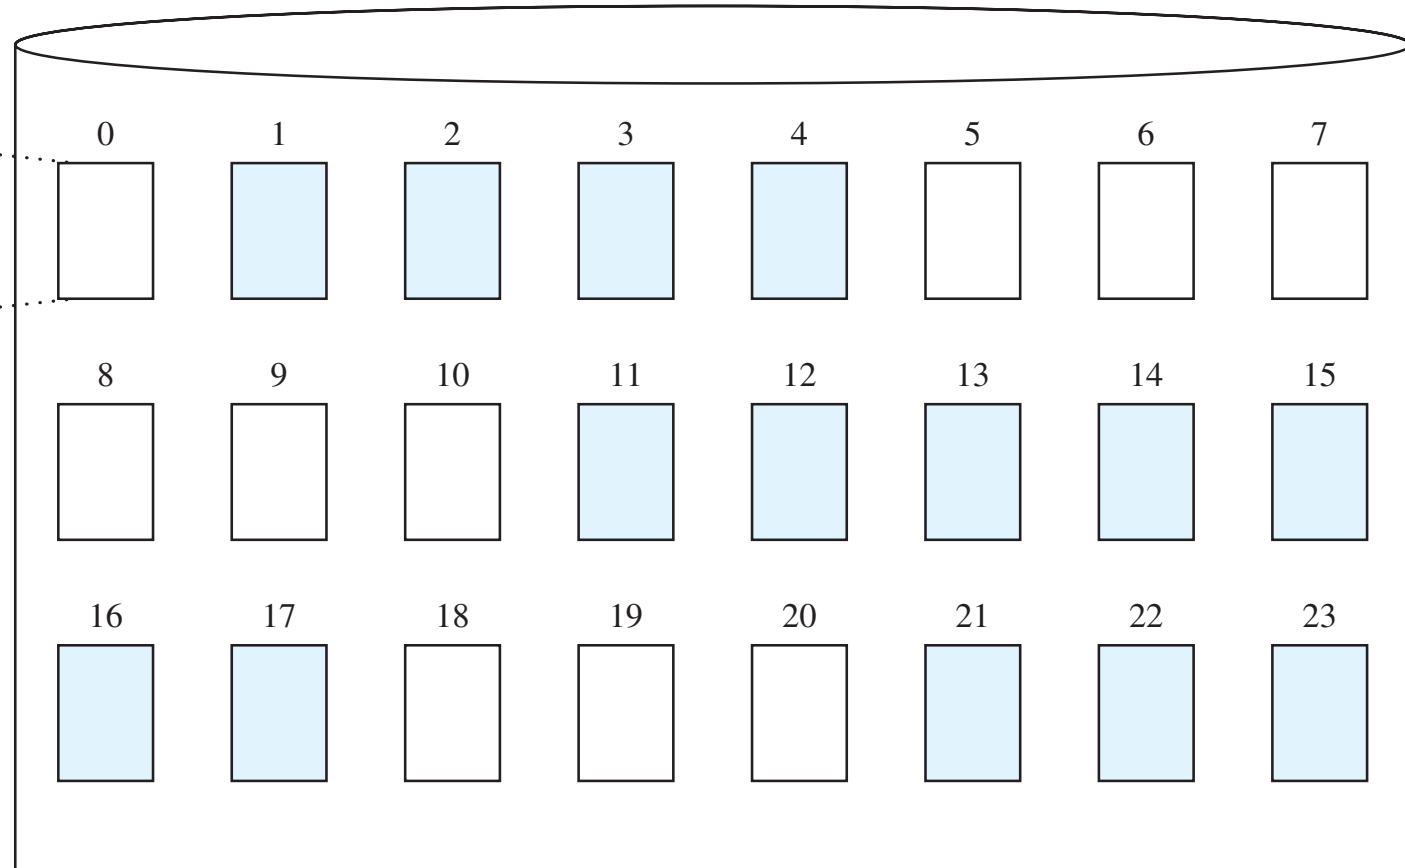
# File allocation techniques

- Contiguous
- Linked
- Indexed


## Contiguous

Directory

<u>File</u>	<u>Start</u>	<u>Size</u>
TermPaper	5	6
MailingList	18	3



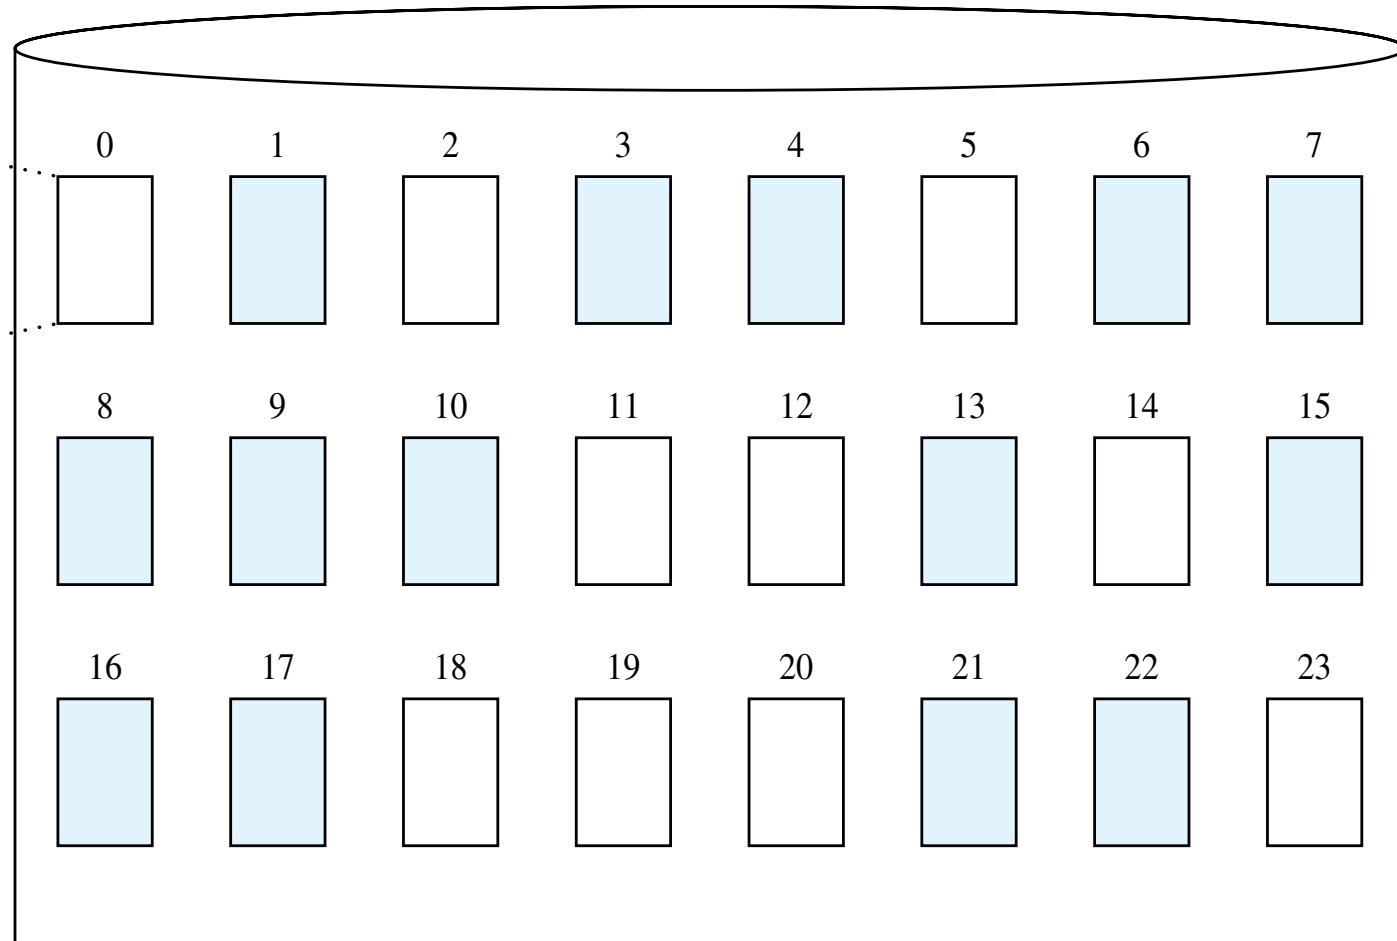
## Linked



## Indexed

### Directory

<u>File</u>	<u>Blocks</u>
TermPaper	5, 11, 12, 14, 23, 20
MailingList	18, 19, 2



# Physical errors

- Parity bits are redundant bits transmitted in addition to the data bits
- Two approaches to the error problem
  - ▶ Detect the error and retransmit or discard the received message
  - ▶ Correct the error



# Codes

- Code
  - ▶ The set of data plus parity bit patterns that are sent
- Code word
  - ▶ An individual pattern from the code


# Code requirements

- The *Hamming distance* between two code words of the same length is the number of positions in which the bits differ
- The *code distance* is the minimum of the Hamming distance between all possible pairs of code words in the code


# Error detection

- To detect  $d$  errors the code distance must satisfy the equation

$$\text{code distance} = d + 1$$



(a) Single-error detecting.




(b) Double-error detecting.


# Error correction

- To correct  $d$  errors the code distance must satisfy the equation

$$\text{code distance} = 2d + 1$$




(a) Single-error correcting.



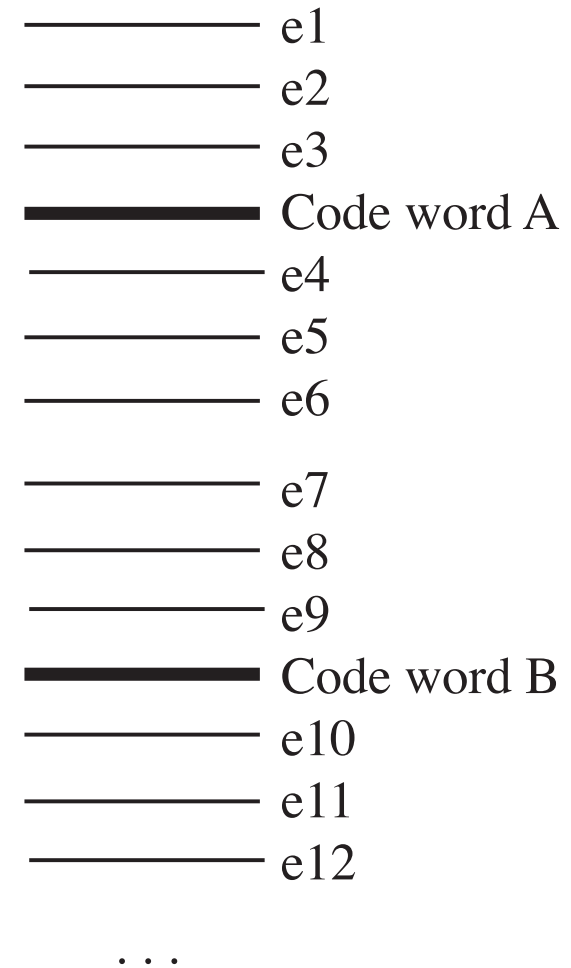
(b) Double-error correcting.

# Single-error-correcting codes

- $m$  data bits
- $r$  parity bits
- $n = m + r$  total number of bits



(a) Code word structure.



(b) Grouping of received words with zero or one error.



# Code requirement

- $m + r + 1 \leq 2^r$
- A *perfect code* is a code for which  
 $m + r + 1 = 2^r$

Data bits <i>m</i>	Parity Bits <i>r</i>	Percent Overhead
4	3	75
8	4	50
16	5	31
32	6	19
64	7	11
128	8	6

# How to determine the parity bits

- Do not append the parity bits at the end of the code word
- Instead, distribute them throughout the code word
- Advantage: A single error can be corrected without having to compare the received word with every code word

1	2	3	4	5	6	7	8	9	10	11	12
		1		0	0	1		1	1	0	0

$$1 = 1$$

$$5 = 1 + 4$$

$$9 = 1 + 8$$

$$2 = 2$$

$$6 = 2 + 4$$

$$10 = 2 + 8$$

$$3 = 1 + 2$$

$$7 = 1 + 2 + 4$$

$$11 = 1 + 2 + 8$$

$$4 = 4$$

$$8 = 8$$

$$12 = 4 + 8$$

Parity bit 1 checks 1, 3, 5, 7, 9, 11

Parity bit 2 checks 2, 3, 6, 7, 10, 11

Parity bit 4 checks 4, 5, 6, 7, 12

Parity bit 8 checks 8, 9, 10, 11, 12

# RAID

- Redundant Array of Inexpensive Disks
- RAID Levels

0: Nonredundant striped

1: Mirrored

01, 10: Striped and mirrored

2: Memory style ECC


3: Bit-interleaved parity

4: Block-interleaved parity

5: Block-interleaved distributed parity

# RAID Level 0



- Nonredundant striped
- Not true RAID because there is no redundancy, therefore no enhanced reliability
- Advantage: Increased performance





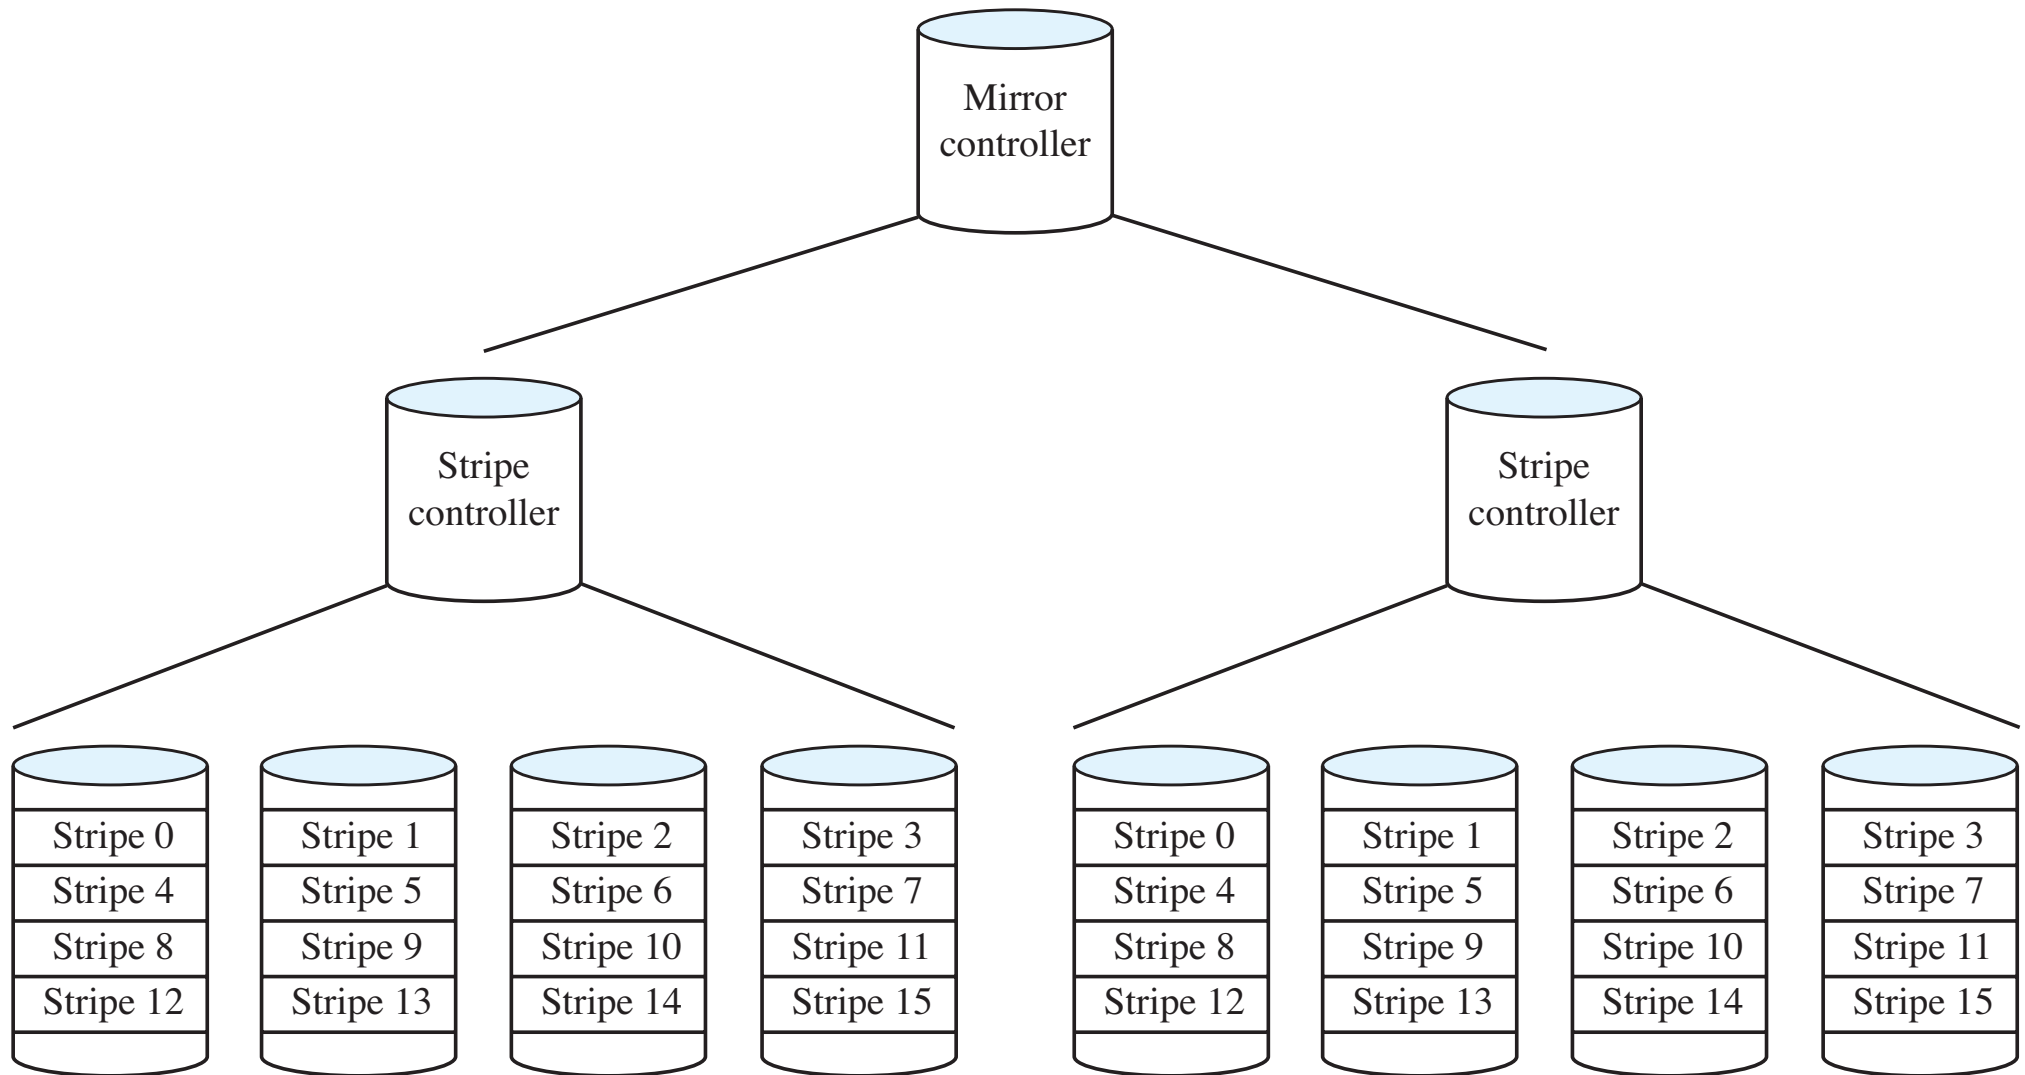
# RAID Level I

- Level I: Mirrored
- An exact mirror image of each disk on a separate drive
- Disk writes are done in parallel to each drive
- Advantage: Reliability
- No performance hit, but no increase in performance either

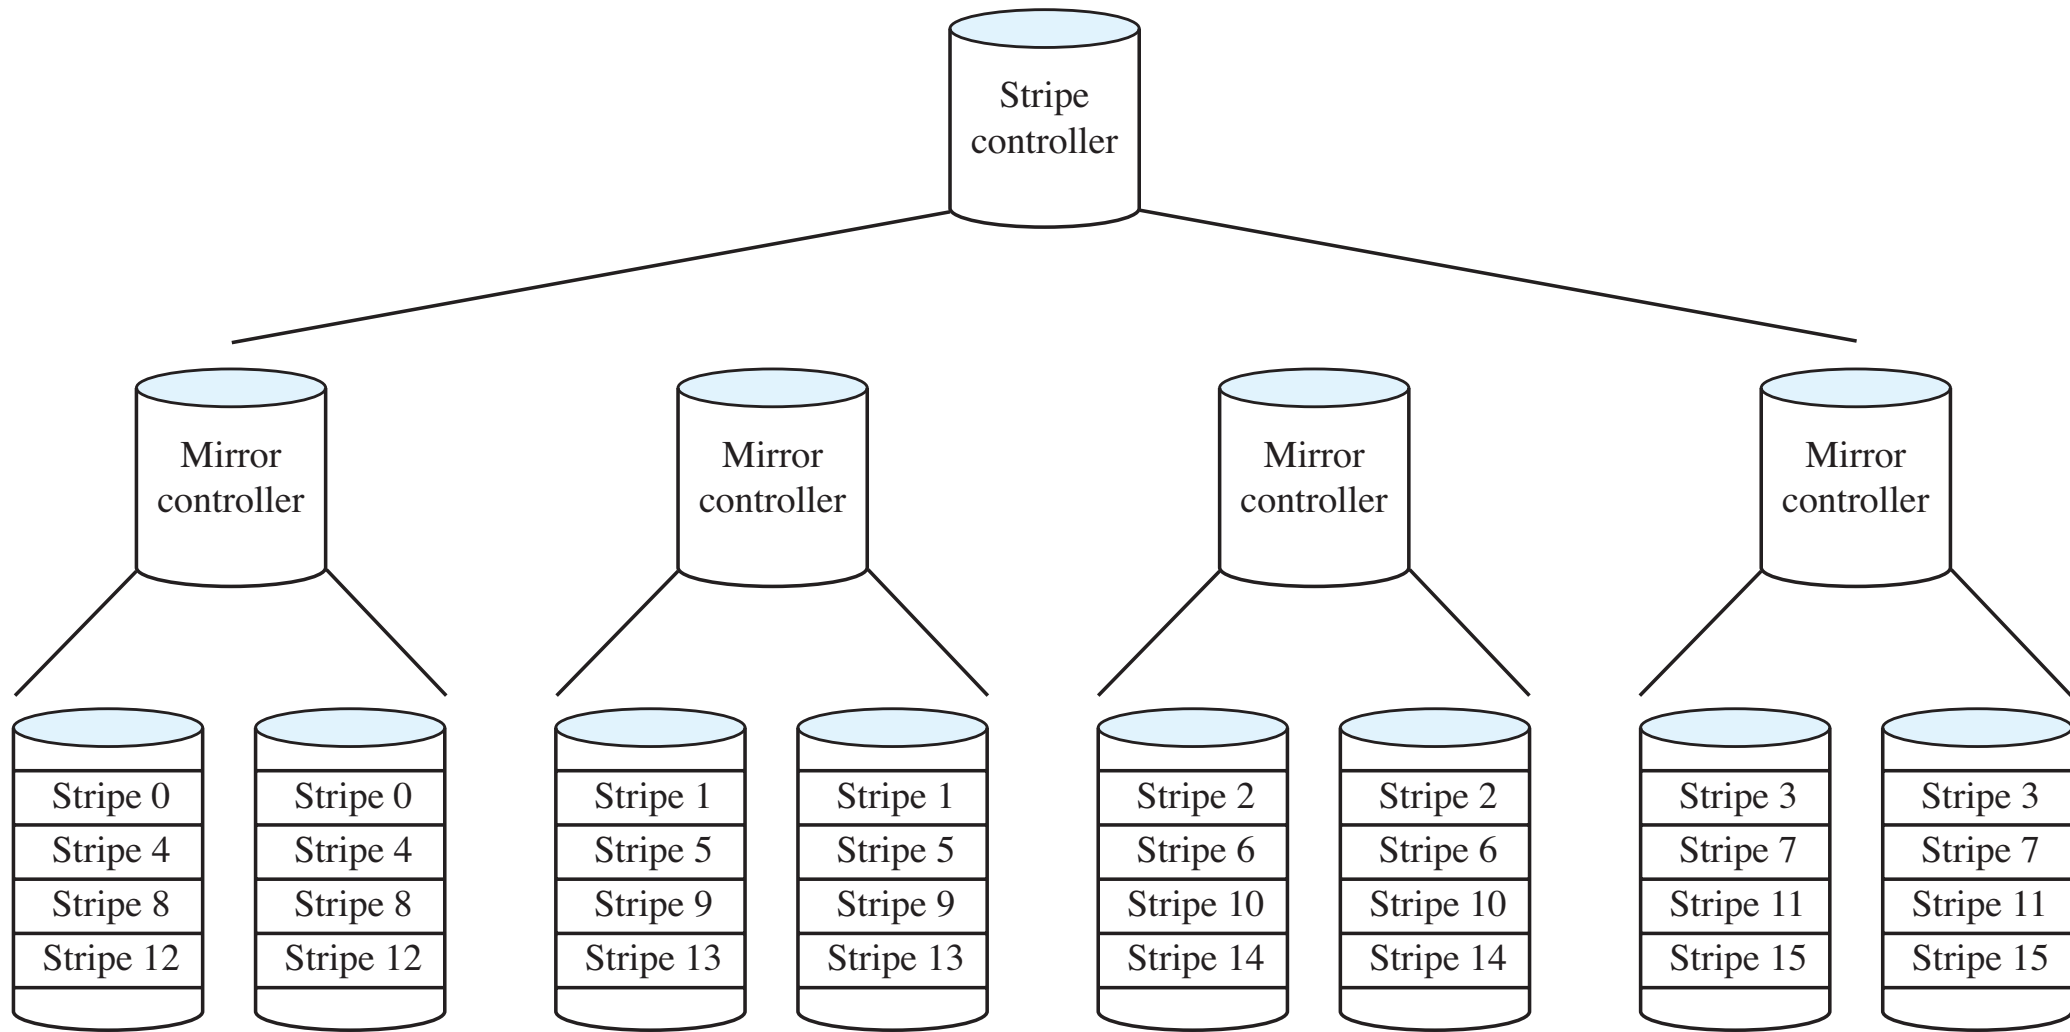


# RAID Levels 01, 10

- Level 01: Mirrored stripes
- Level 10: Striped mirrors
- Both increased reliability and increased performance
- Advantage: Easy to hot-swap an entire drive
- Disadvantage: Huge storage overhead. Not as efficient as using parity to correct errors




(a) RAID Level 01, mirrored stripes.



(b) RAID Level 10, striped mirros.

# RAID Level 2


- Memory-style ECC
- Error correction at the bit level
- Disks must be synchronized
- Used on older supercomputers with 32 data bits and 6 parity bits
- Modern disks have internal bit-level ECC, so Level 2 is no longer used commercially



# RAID Level 3

- Bit-interleaved parity
- When a drive fails, you know which disk drive it is
- Knowing which bit fails, it takes only *one* parity bit to *correct* the error
- More efficient use of disk space





# RAID Level 3 disadvantages

- To replace a failed drive, you must read *all* the drives to correct the error and compute the data in the substitute drive
- Every read/write request requires you to access every drive
- No concurrency, so no performance benefit

# RAID Level 4

- Block-interleaved parity

- Example

Stripe 0: Bits 0 through 1023

Stripe 1: Bits 1024 through 2047


Stripe 2: Bits 2048 through 3071

Stripe 3: Bits 3072 through 4095

Bit 1 of P0-3 checks 0, 1024, 2048, 3072

Bit 2 of P0-3 checks 1, 1025, 2049, 3073

...



# Level 4 advantages

- Striping is not at the bit level, so rotational synchronization is not required
- Better performance than Level 3 for small random read requests
  - ▶ With Level 4, many small files on different disks allow for concurrency
  - ▶ With Level 3, many small files must be read sequentially

# Level 4 disk write without shortcut

- To write to stripe 0:
  - ▶ Read stripes 1, 2, 3
  - ▶ Compute parity with stripe 0
  - ▶ Write to stripe 0 and P0-3
- Result is *three* reads and *two* writes

# Level 4 disk write with shortcut

- To write to stripe 0:
  - ▶ Read old stripe 0 and old stripe P0-3
  - ▶ If the new stripe 0 bit differs from the old stripe 0 bit, flip the corresponding bit in P0-3
  - ▶ Write to stripe 0 and P0-3
- Result is *two* reads and *two* writes

# Level 4 disadvantage

- With or without the shortcut, every write request must write to the parity disk
- The parity disk becomes the performance bottleneck



# RAID Level 5

- Block-interleaved distributed parity
- Rather than store all the parity on one disk, the parity information is scattered among all the disks
- No one disk has the responsibility for the parity information of the whole array
- Better performance with reliability maintained

